



State-of-the-art App Containers

Andrej Galad, Shivam Maharshi



Hypothesis

Given the current state of the two application containers, we believe that:


- Docker is a more logical choice over rkt from the aspect of performance, features and container clustering support.
- In a long term rkt has a potential to catch up with/succeed Docker.

Evaluation Points


| # | Criteria | Notes |
|----|-----------------------------------|---|
| 1. | Ideology | Motivation, specification standardization and security concerns |
| 2. | Architecture | Terminology and architecture |
| 3. | Lifecycle | Startup & Lifecycle of a container |
| 4. | OS Support | OS & package manager support |
| 5. | Feature Support | Network Types, Volume Mounting, Image Building, Registry Availability, Dynamic Image Creation, Image Distribution |
| 6. | Performance Benchmarks | Startup Time, CPU Compute, Network Performance, File IO Bandwidth, Distributed Processing Benchmark, Load Benchmark |
| 7. | Clustering | Turn key cloud solutions & Kubernetes |
| 8. | Community + Documentation Support | Community size, development bug resolving pace and documentation |

Container Overview


- Ideology
- Lifecycle
- Container Management
- Image Distribution



Ideology



- rkt
 - Standard Container - App Container Specification
 - Improved Security - Privilege Separation
 - Easy multiple implementations possible
 - Unix Philosophy - Components Isolation & Clear Integration Points
- Docker
 - Vendor-specific interface
 - Push towards platform - Machine, Compose, Swarm, Cloud...
 - Tightly coupled components - Docker Engine ↔ Containers



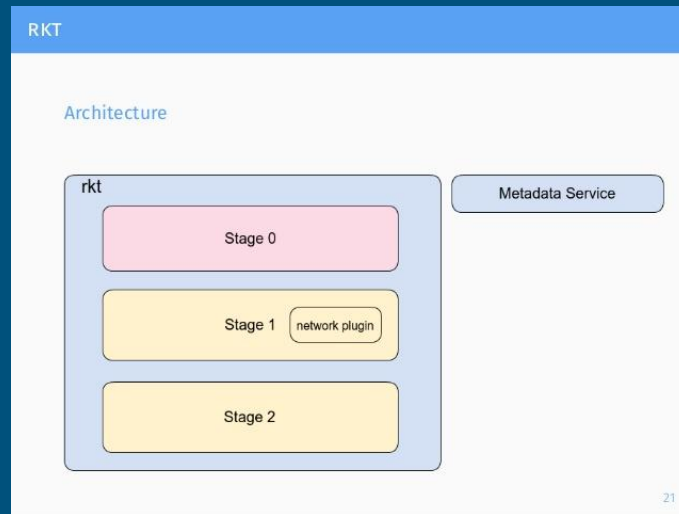
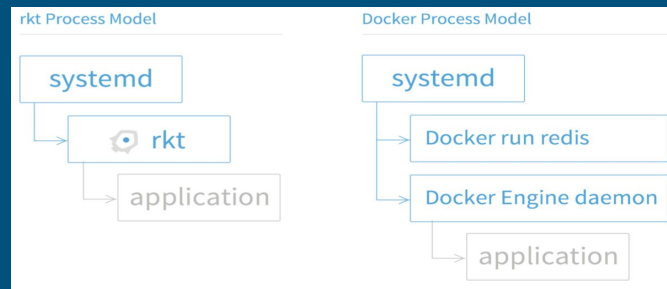
Life Cycle & Startup

- **rkt**

- 4 Phases of container - Prepare, Run, Exited Garbage & Garbage
- Execution in 3 steps - (Hence the pluggable isolation)
 - Stage0 (stage1 aci, generate manifest, uuid)
 - Stage1 (create cgroups, namespaces pod)
 - Stage2 (runs the fs prepared)
- Linear execution

- **Docker**

- Blackbox - Docker Engine
- Steps
 - Fetch Image from Docker registry
 - Setup file system with read-write layer
 - Setup networking for docker-host communication
 - IP address is attached to running container
- Circular execution



Architecture

- **rkt**

- Configures AC namespaces using systemd-nspawn
- Pods - Basic units of execution (Kubernetes)
- Apps in a Pod share context
- Can be updated in-place

- **Docker**

- Client-Server model
 - Docker Engine + CTL
- Containers - orchestrated by daemon
- Basic unit of execution is an AC

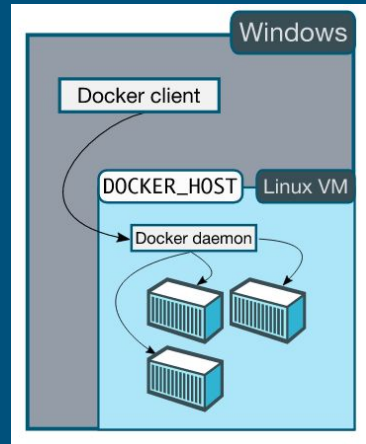
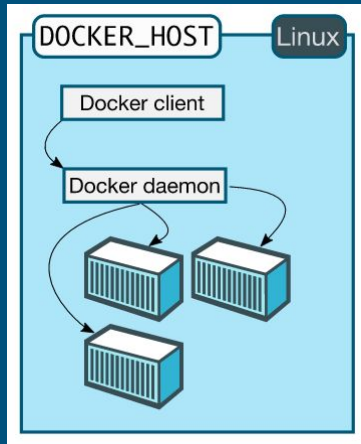
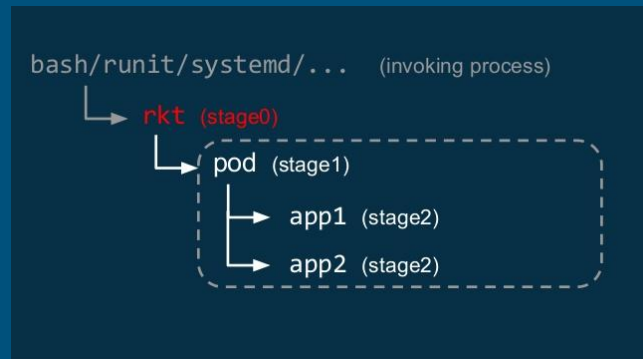
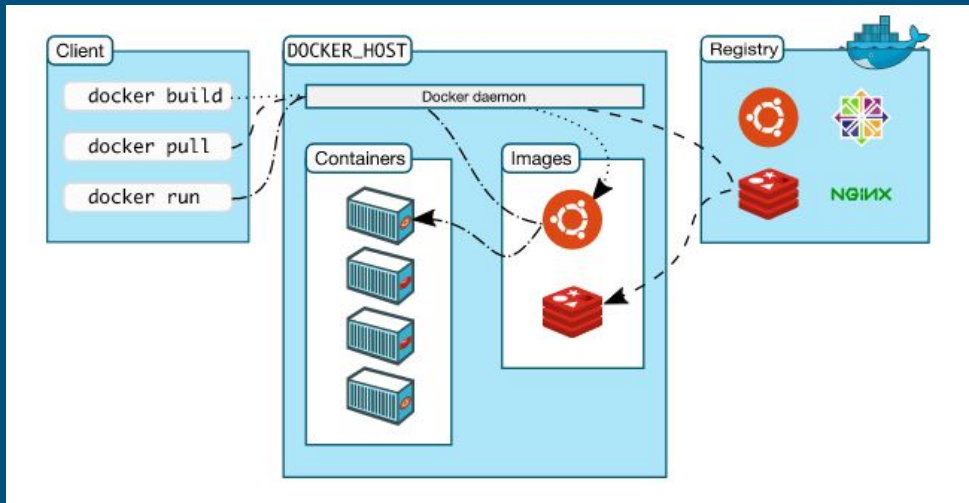


Image Distribution

- Rkt (distributed)
 - Plain tarball images over HTTP
 - DNS discovery of custom namespace & signatures
 - docker2aci
- Docker (centralized)
 - Primary - Docker Hub
 - Secondary - 3rd party registries
 - Cloud providers + private repos



Performance, performance, performance

- Startup
- CPU
 - Sysbench, CoreMark
- Network
 - iPerf3
- File I/O
 - fio
- Distributed processing
 - Fedora 4 Benchmark
- Load
 - Wikipedia Benchmark




Performance, performance, performance

- Startup
- CPU
 - Sysbench, CoreMark
- Network
 - iPerf3
- File I/O
 - fio
- Distributed processing
 - Fedora 4 Benchmark
- Load
 - Wikipedia Benchmark



Startup Time

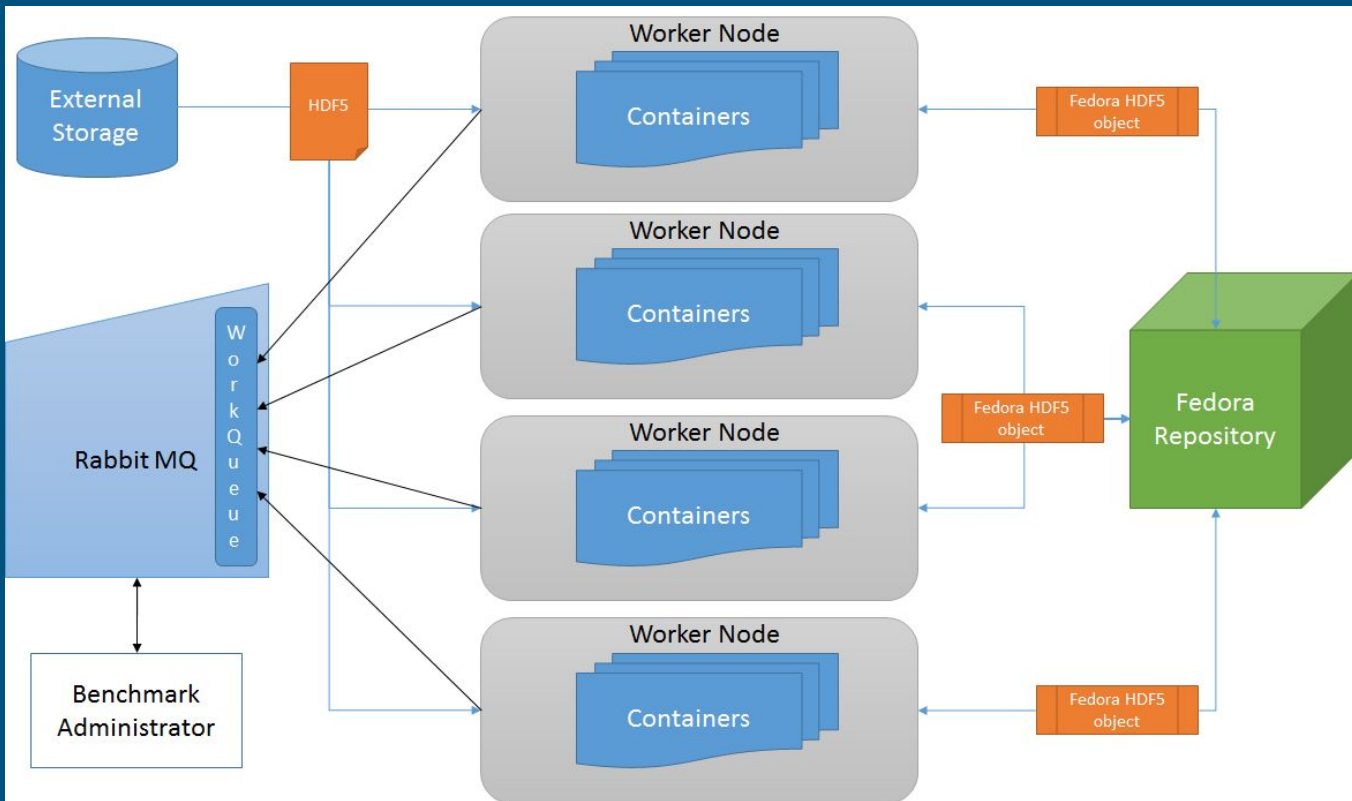
- “Vanilla” machine
 - Docker 1.11.1
 - Rkt 1.3.0 -> 1.5.0
- Image - busybox:latest
- Sequential startup of 20 containers



Distributed Processing Benchmark

- Fedora 4 Repository - <http://fedorarepository.org/>
 - ModeShape/Infinispan
- Input
 - Goodwin Hall sensor data - HDF5
 - 180 files of approx. 50MB -> 9GB (subset of 2 day data collection)
- Chameleon - <https://www.chameleoncloud.org/>
 - OpenStack KVM
 - Bare Metal
- Fedora Benchmark - <https://github.com/VTUL/VT-Fedora-Benchmark>
 - 3 workflows - Ingestion, Fixity checking, FFT

Fedora Benchmark Overview


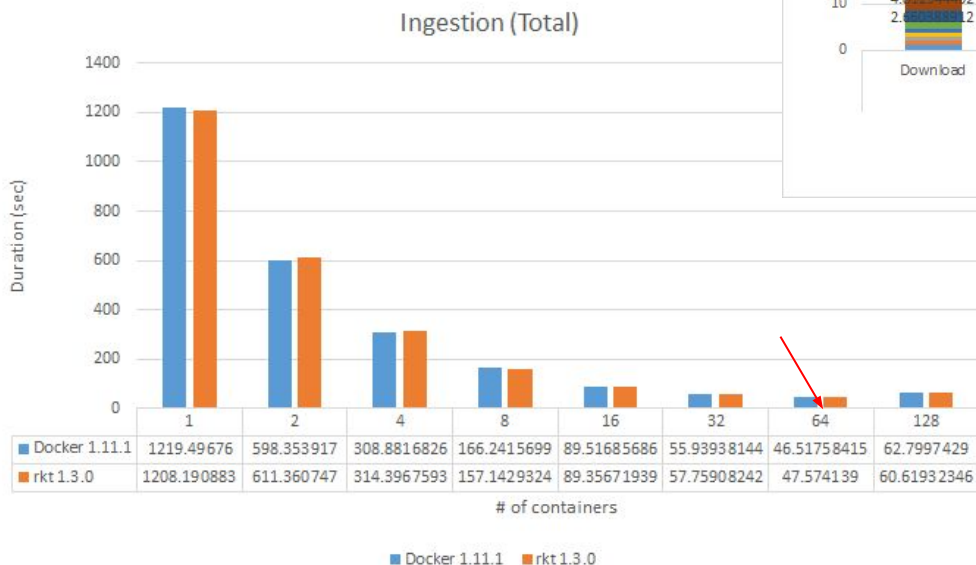


Environment Specs + Setup

- Chameleon OpenStack KVM (Ubuntu Server 14.04 LTS)
 - RabbitMQ
 - m1.medium, 2 VCPUs, 4GB RAM
 - Fedora 4
 - m1.large, 4 VCPUs, 8GB RAM, 80GB drive
- Chameleon Bare Metal (Ubuntu 14.04)
 - 4 Worker Nodes
 - 48 VCPUs, 128GB RAM, 230 GB drive

| Phase | # of hosts/# of containers | Total |
|-------|----------------------------|-------|
| 1 | 1 host with 1 worker | 1 |
| 2 | 2 hosts with 1 worker | 2 |
| 3 | 2 hosts with 2 workers | 4 |
| 4 | 4 hosts with 2 workers | 8 |
| 5 | 4 hosts with 4 workers | 16 |
| 6 | 4 hosts with 8 workers | 32 |
| 7 | 4 hosts with 16 workers | 64 |
| 8 | 4 hosts with 32 workers | 128 |



Ingestion



Steps:

1. Download HDF5 file from external storage - Google Drive
2. Extract metadata/headers using FITS
3. Create Fedora object and ingest data



Fixity Checking



Steps:

1. Fetch HDF5 checksum from Fedora
2. Reconstitute Fedora object and generate checksum
3. Compare and store result as object metadata

Fast Fourier Transform



Steps:

1. Reconstitute Fedora object
2. Compare FFT (numpy.fft - $n \cdot \log(n)$)

2nd run - Setup

- Bring everything together...
- Chameleon Bare Metal
 - 48 VCPUs, 128GB RAM, 230 GB drive
- Fedora 4 + Benchmark + RabbitMQ
 - No network isolation

| Phase | # of containers |
|-------|-----------------|
| 1 | 1 |
| 2 | 2 |
| 3 | 4 |
| 4 | 8 |
| 5 | 16 |
| 6 | 32 |
| 7 | 64 |
| 8 | 128 |

Distributed Load Benchmark - Server Side

- Two configurations
 - Single Node Multiple Containers
 - Multiple Node Two Containers
- Wikipedia Server
 - Realistic Application
 - Realistic Workload
- Yahoo Cloud Service Benchmark
 - Custom Web Module
 - Link: <https://github.com/shivam-maharshi/YCSB4WebServices>

Single Node Multiple Containers

- YCSB

- OS X El Capitan
- 4 x 2.66 GHz Intel Core i5
- 4 GB 1.06 GHz DDR3

- Nginx Load Balancer

- OS X El Capitan
- 8 x 3.1 GHz Intel Core i5
- 32 GB 1.3 GHz DDR3

- Node

- Ubuntu 14.04
- 8 x 2 GHz Intel Core i5
- 16 GB

- Version

- Apache 2.2, MediaWiki 1.26.2, PHP 5.5, MySQL 5.x



Multiple Node Two Containers

- YCSB

- OS X El Capitan
- 4 x 2.66 GHz Intel Core i5
- 4 GB 1.06 GHz DDR3

- Nginx Load Balancer

- OS X El Capitan
- 8 x 3.1 GHz Intel Core i5
- 32 GB 1.3 GHz DDR3

- Node #

- Ubuntu 14.04
- 2 x 2 GHz Intel Core i5
- 4 GB

- Version

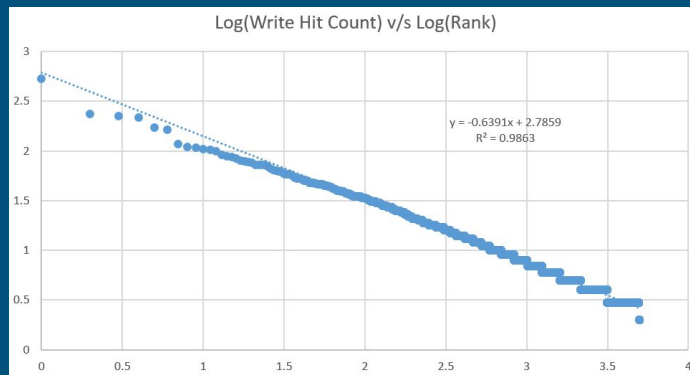

- Apache 2.2, MediaWiki 1.26.2, PHP 5.5, MySQL 5.x



Wikipedia Workload

- Trace Generation

- Greek - 400k pages - 4GBs
- 1 Month Trace - Jan 2016
- Read = 2.8 Billion
- Write = 60 K
- Sort URLs in Descending Hits
- Read & Write - Zipf's Distribution
- Read Zipf's Constant = 0.6175
- Top 10k pages, $R^2=0.9898$
- Write Zipf's Constant = 0.6391
- Top 5k pages, $R^2=0.9772$





Benchmarking



- Yahoo Cloud Services Benchmark
 - Read & Write Trace
 - Parameters - Zipf's K, R/W Ratio..
 - Execution Timeout = 10 secs
- Step Benchmarking
 - Incremental Workload - In steps of 25 RPS
 - Sufficient Duration - 2 Hours / Conf
 - Resource Monitoring - CPU, RAM & Network
- Metrics
 - Latency
 - Throughput
- Bottleneck
 - CPU Bound



Single Node Multiple Containers



Multiple Node Two Containers



Cluster it up...



- Turn-key Cloud Solutions - rkt though luck :(
 - EC2 Container Service - custom solution
 - Azure Container Service - Swarm (NO LONGER PREVIEW :D)
 - Google Container Engine - Kubernetes (gcloud, kubectl)
- Kubernetes
 - Docker - main runtime (even for CoreOs)
 - Rkt - can be configured as runtime but...
 - Only nodes
 - No automation
 - <http://146.148.35.100:8080/>
- Rktnetes - rkt 1.7.0

Conclusion

- Docker - still no.1 container on the market
 - Deployment/provisioning, OS availability, startup time, features, 3rd party support, platform...
 - Both developers and ops
- Rkt's main potential - simplicity and modularity
 - Emphasis on customizability
 - Adherence to a common standard
 - "Made for Kubernetes"